SAP HANA Overview

Lesson Objectives

- **On Completing this course, participants will be able to:**

  - Understand SAP HANA
  - Understand ABAP ON HANA
  - Understand SAP HANA System Architecture
  - To know SAP IN-Memory Strategy and Technology
  - Understand Row Store and Column Store

Contents

- SAP HANA System Landscape Connectivity Overview
- Technology Innovations and HANA Architecture
- In-Memory Architecture
- Parallel Processing
- Row Store, Column Store and Dictionary Compression
- ABAP Development shortcuts
- SAP HANA Coding Pattern

Contents

- Introduction to SAP HANA Studio
- SAP HANA Mandatory Adaptations
- SAP HANA Coding Pattern
- ABAP Development shortcuts

AB1011 - ABAP ON HANA

HANA – **H**igh-performance **AN**alytic **A**ppliance

HANA – High-performance ANalytic Appliance

SAP Business Suite Powered by SAP HANA

What makes HANA so Unique?

- 100% In-Memory computing (Allow OLTP & OLAP processing in Real-Time) Memory
- 5-20x Compression (Column based compression)
- Up to 10,000x Faster (Massive parallel scaling)

ABAP Platform and SAP HANA

Business values and target groups

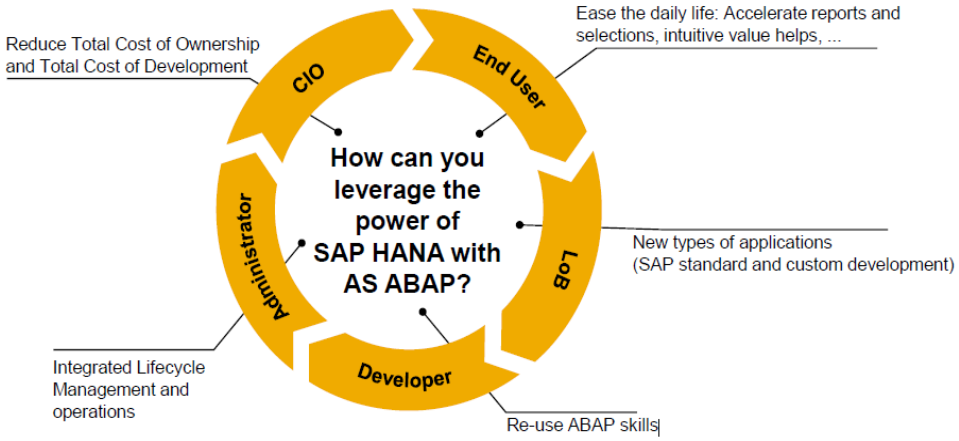

SAP HANA EVOLUTION

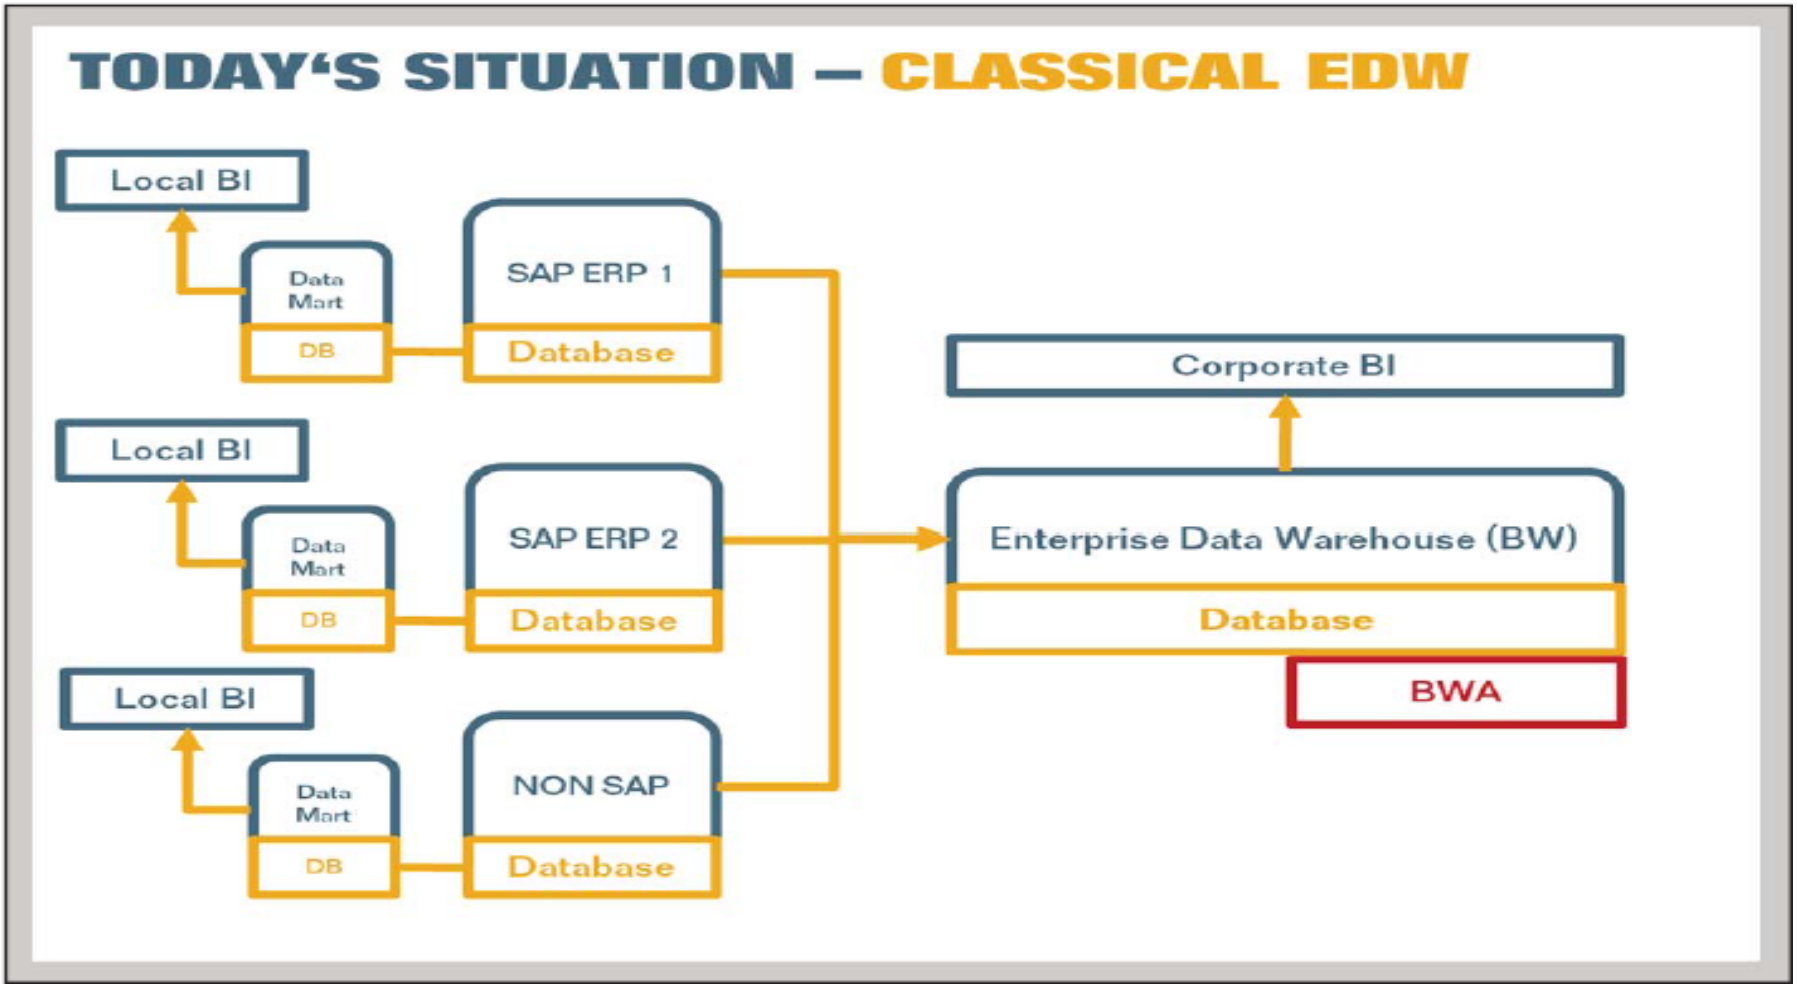

SAP HANA EVOLUTION

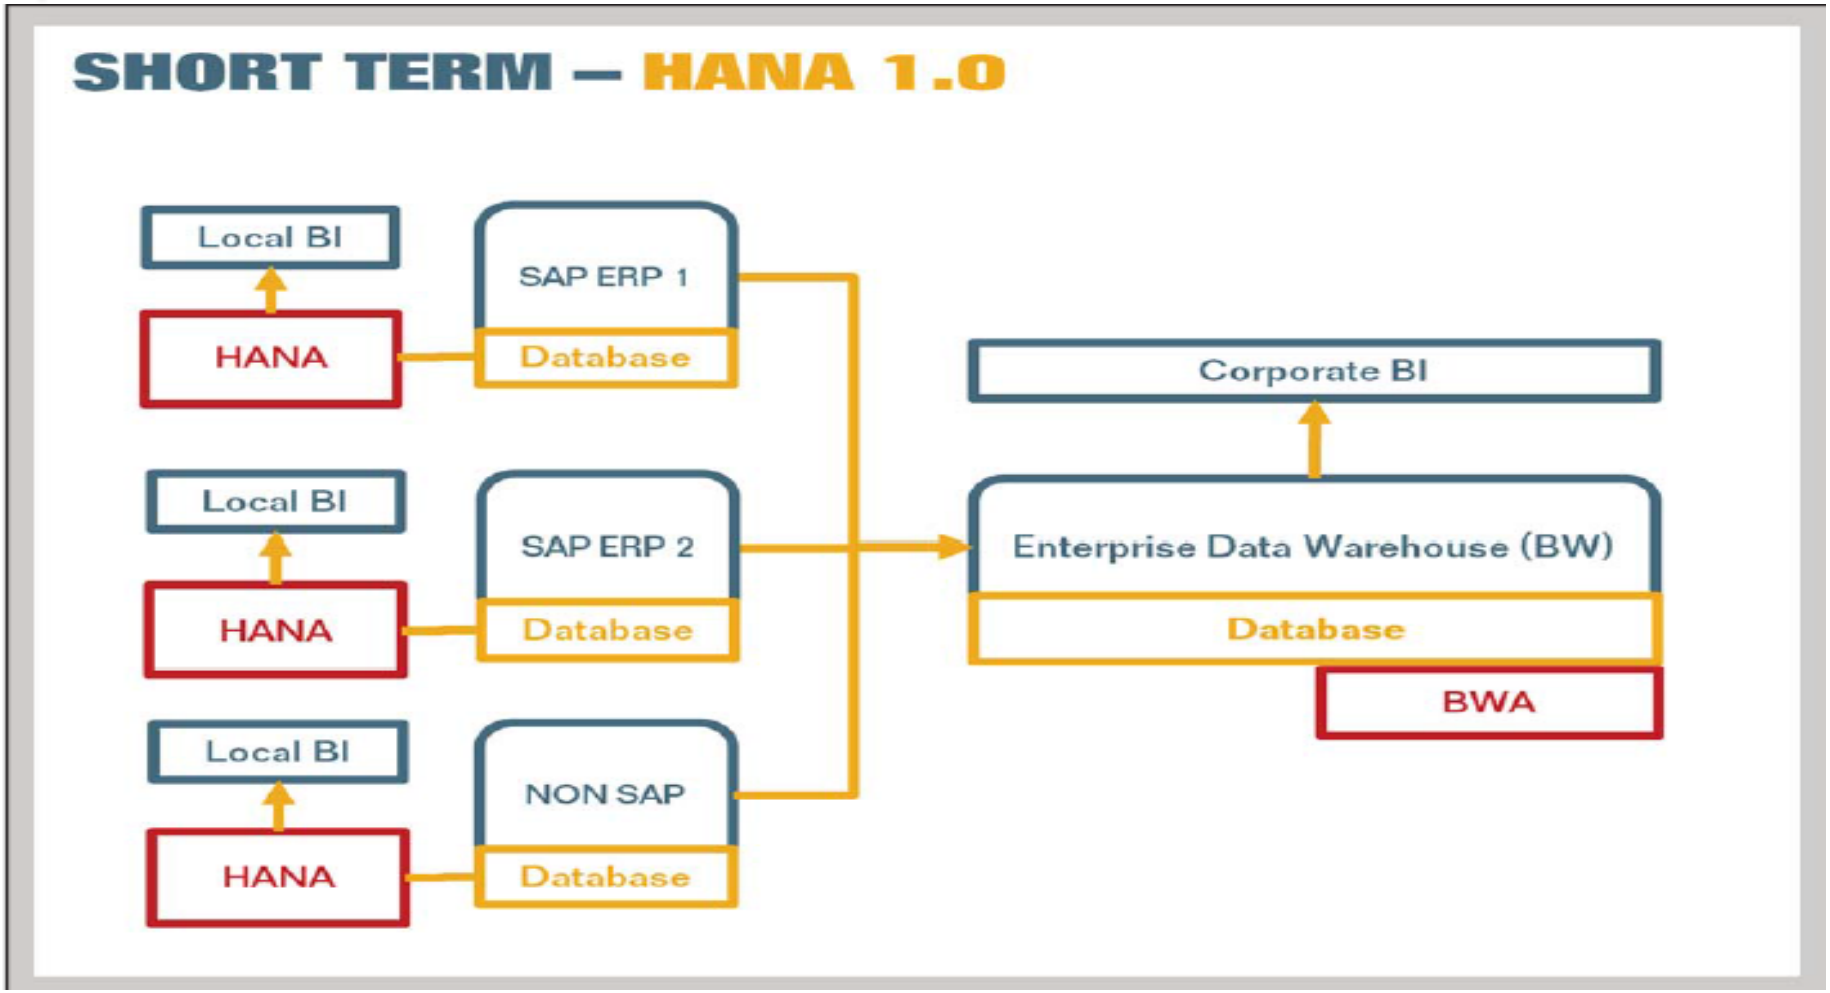

SAP HANA EVOLUTION

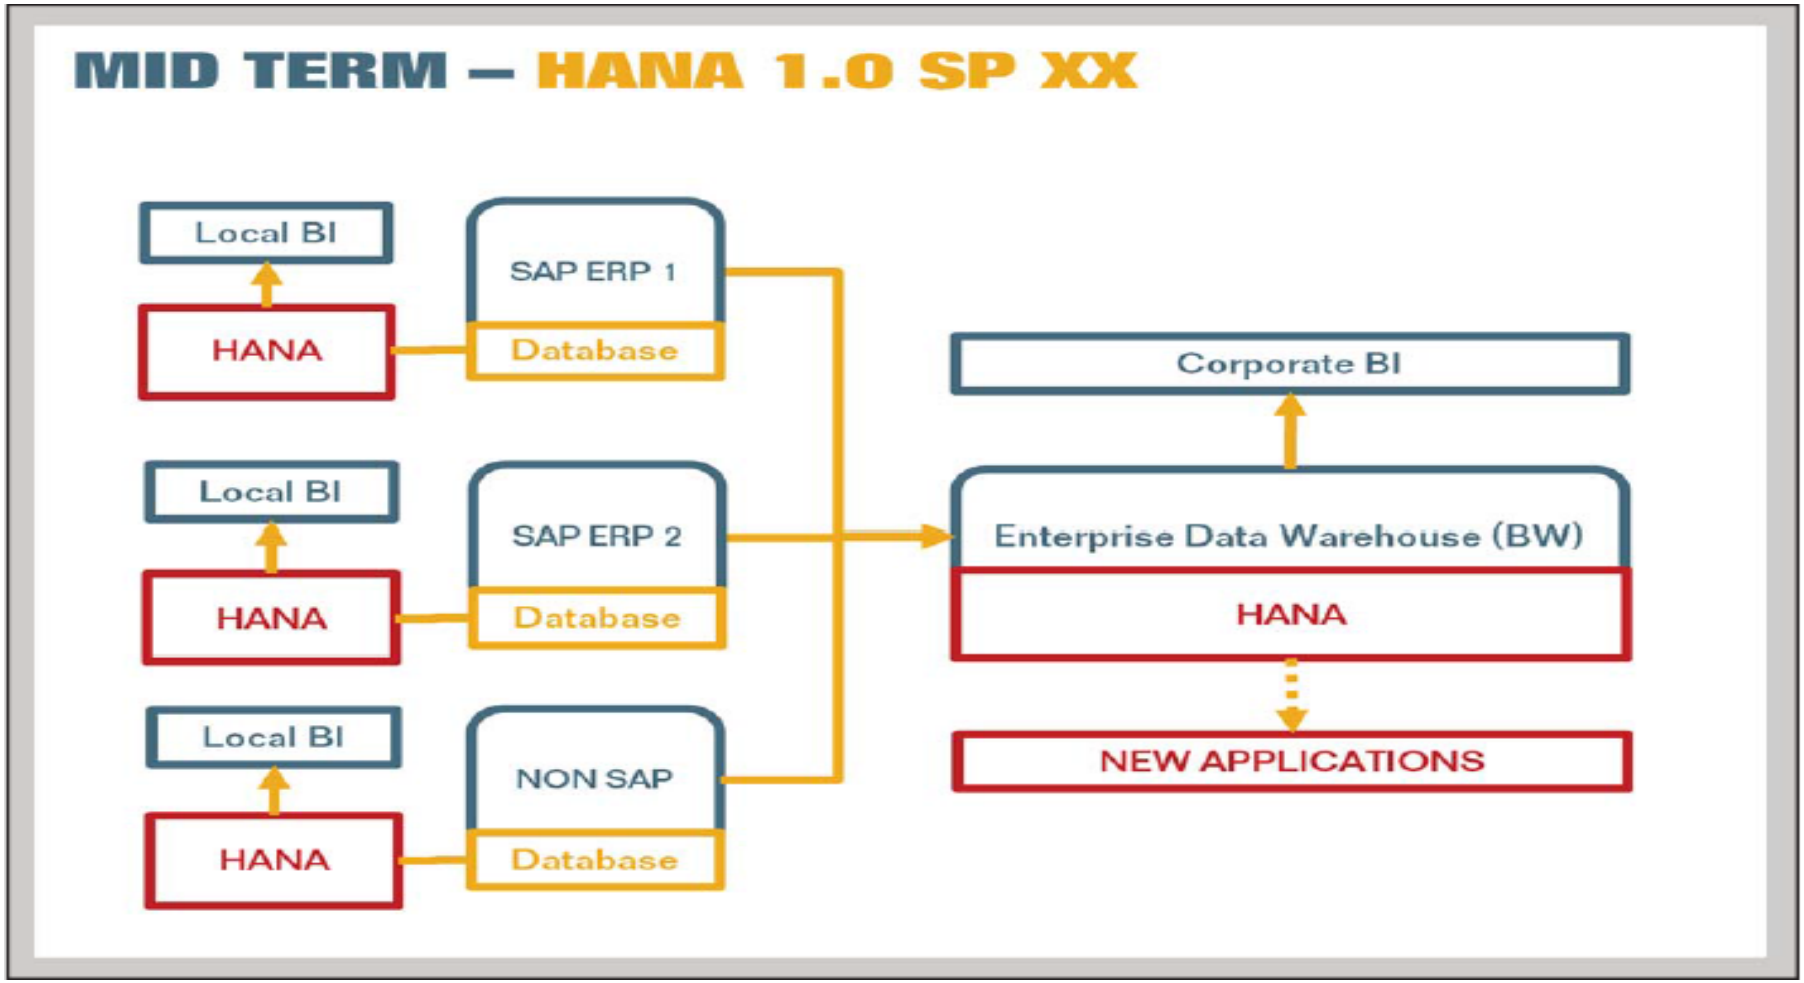

SAP HANA EVOLUTION

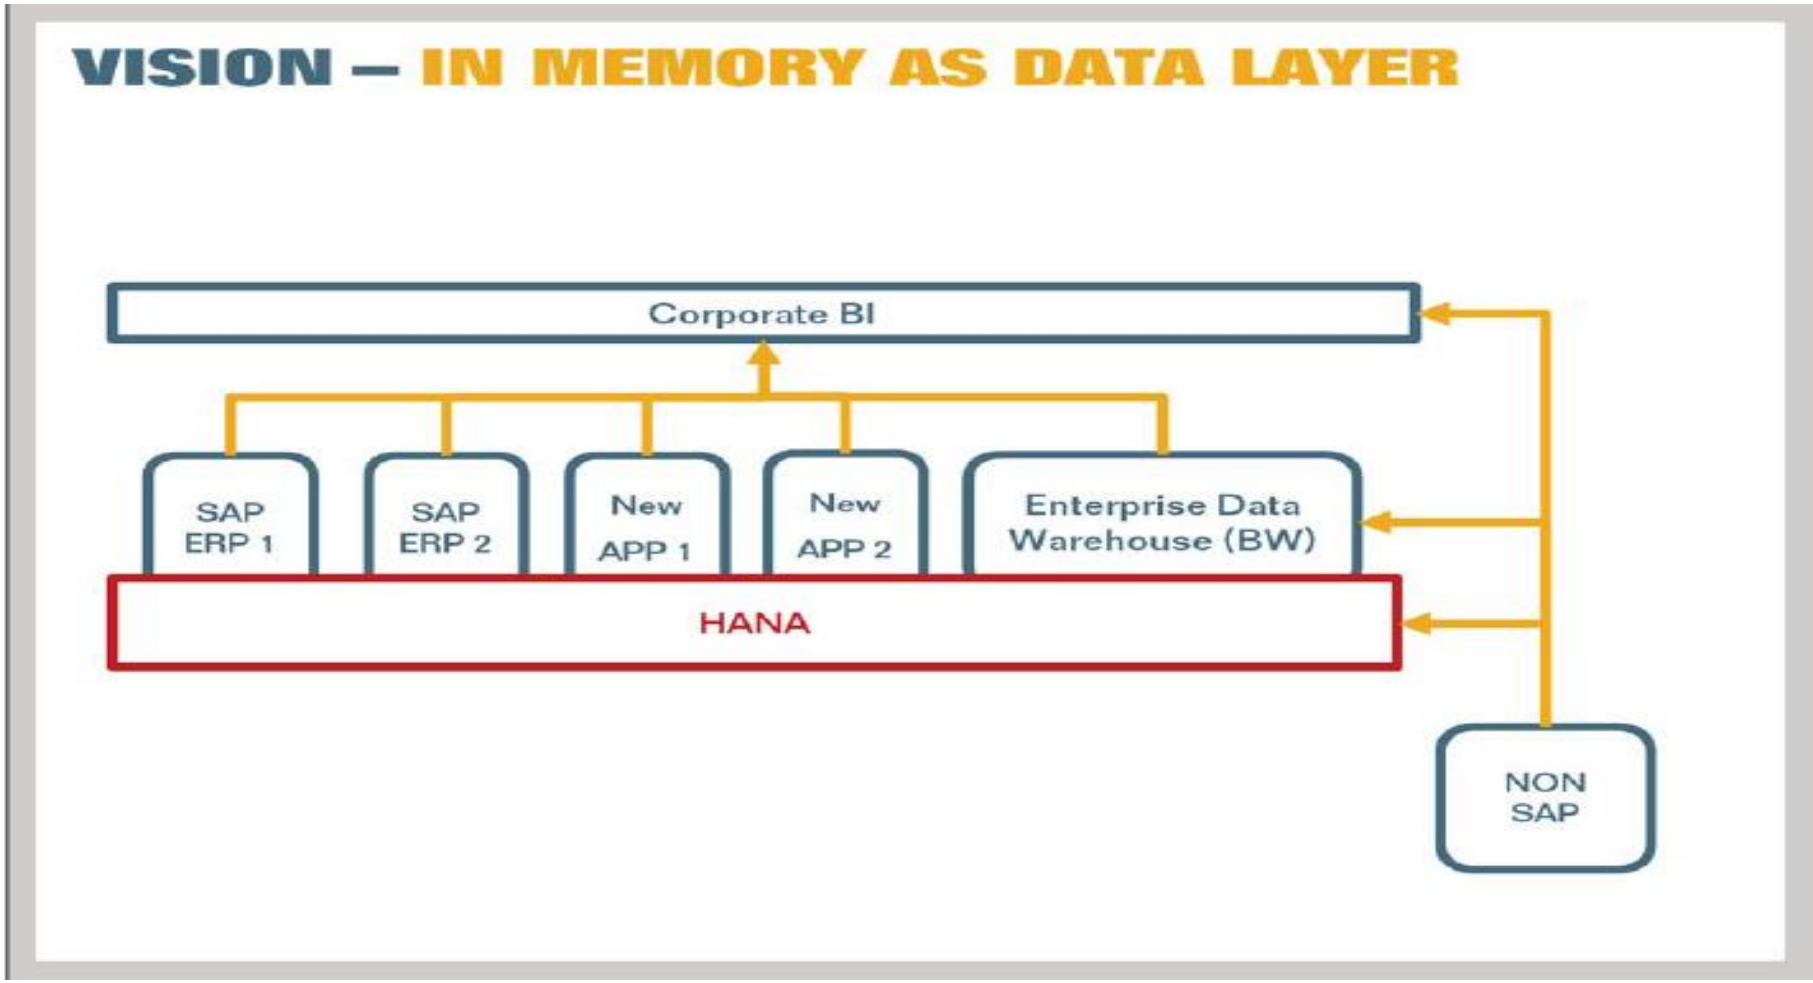

SAP HANA System Landscape Connectivity Overview

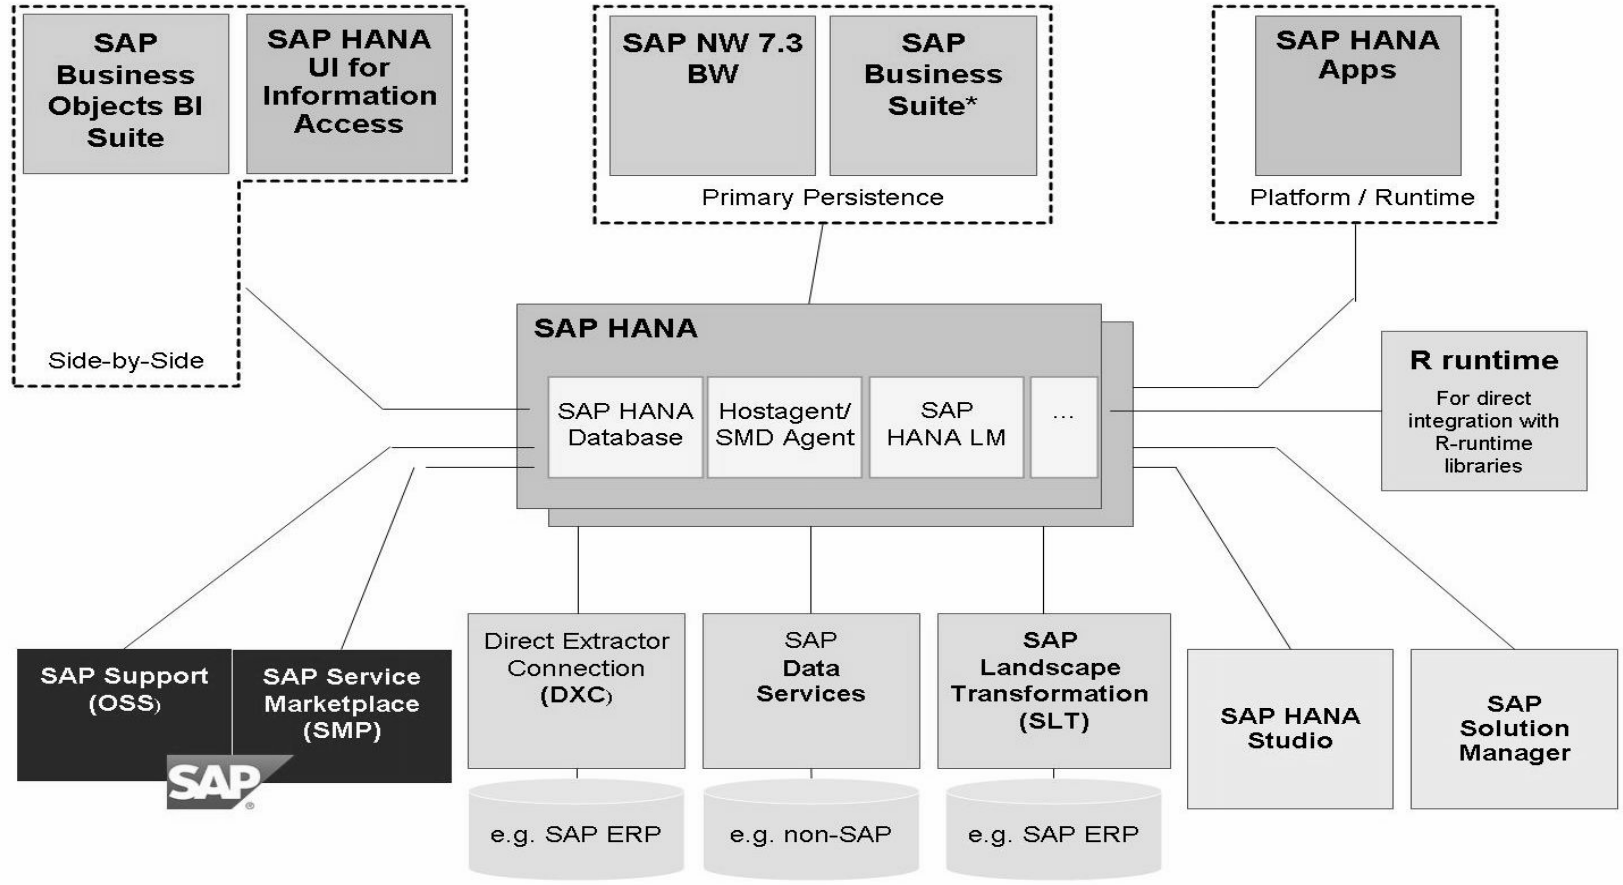

Technology Innovations

Dramatically improved hardware economics and technology innovations in software have made it possible for SAP to deliver on its vision of the Real-Time Enterprise with in-memory business applications.

HW Technology Innovations

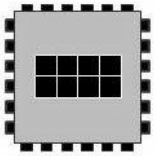

Multi-Core Architecture
(8 CPU x 10 Cores per blade)

Massive parallel scaling with many blades

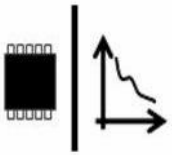

64bit address space – 2TB in current servers

Dramatic decline in price/performance

SAP SW Technology Innovations

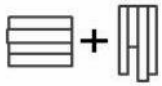

Row and Column Store

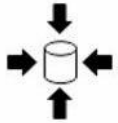

Compression

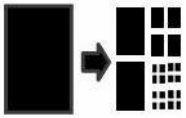

Partitioning

No Aggregate Tables

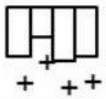

Insert Only on Delta

HANA Architecture

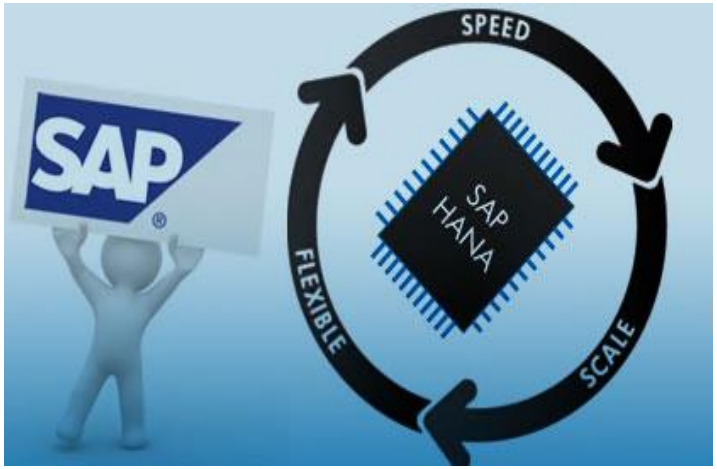

What is SAP HANA?

- In the Conventional Data base, transmitting data from that server to all the application servers and back created a huge performance bottleneck. This is due to I/O operations become very high on multiple APP servers with Single DB.

- SAP HANA combines database, data processing, and application platform capabilities in-memory. The platform provides libraries for predictive, planning, text processing, spatial, and business analytics.

- This new architecture enables converged OLTP and OLAP data processing within a single in-memory column-based data store with ACID compliance, while eliminating data redundancy and latency.

- By providing advanced capabilities, such as predictive text analytics, spatial processing, data virtualization, on the same architecture, it further simplifies application development and processing across big data sources and structures.

- This makes SAP HANA the most suitable platform for building and deploying next-generation, real-time applications and analytics.

The Solution: In-Memory Architecture

- SAP HANA runs on multi-core CPUs with fast communication between processor cores, and containing terabytes of main memory. With SAP HANA, all data is available in main memory, which avoids the performance penalty of disk I/O. Either disk or solid-state drives are still required for permanent persistency in the event of a power failure or some other catastrophe. This does not slow down performance, however, because the required backup operations to disk can take place asynchronously as a background task.

Multicore CPUs
10 Cores / CPU

Multi-CPU Boards
8 CUPs / Board

Multi Server Board
x Boards

Massive Memory setups
2 TB/Server

320 CORES and more!
4 TB RAM and more!

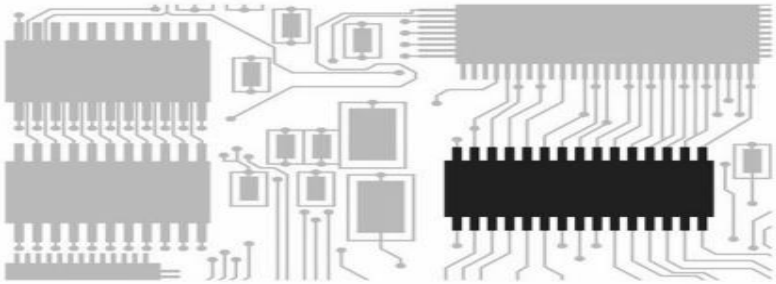

Architecture Overview

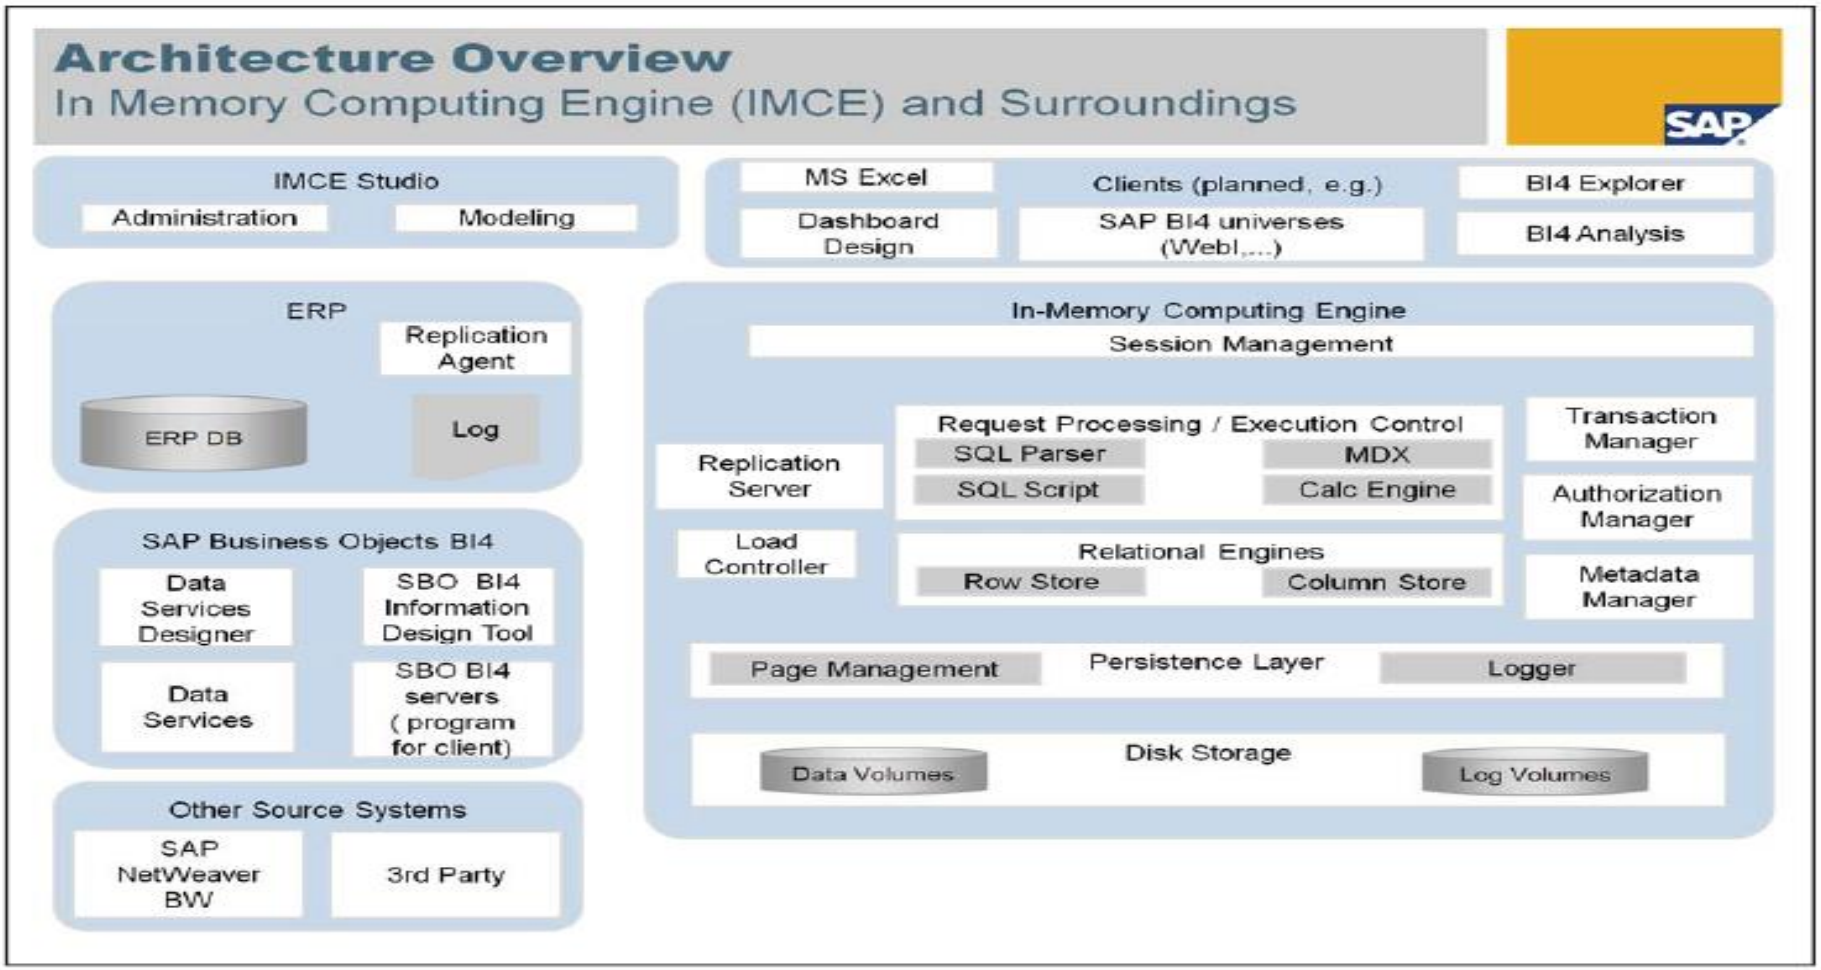

Architecture Overview

- **Session Management:** Connection and Session Management which creates and manages sessions and connections for the database clients. For each session a set of parameters is maintained such as e.g. auto commit settings or the current transaction isolation level.

- **Request Processing and Execution Control:** The client requests are analyzed and executed. Once a session is established, database clients typically use SQL statements to communicate with the in-memory computing engine. For analytical applications the multidimensional query language MDX is supported.

- **Relational Engines:** Are memory based engines in which we store data in Row Store and Column Store. For we need the Persistence layer for data safety to overcome the power cut or OS reboot would permanently erase all data which are currently updating in the tables.

Architecture Overview

- **Persistence Layer:** The persistency layer handles page management and logging(redo and Undo logs) and permanently stores data in a disk storage. This storage has separate volumes for data and log.
- **Transaction Management:** In order to provide consistent views of the data at any given point in time (an ongoing transaction must only see that part of the data that was committed before that transaction was started).
- **Administration Manager:** The In-Memory Computing Studio has an administration component.
- Starting/stopping the In-Memory Computing Engine (upon start, the in-memory stores are reconstructed from the persistence layer).
- User administration including creating/deleting users and authorizations.
- Table administration, including creating indexes or some part of the configuration for data replication.
- Creating or replaying a backup.

Architecture Overview

- Replication Server and Load Controller: are the engine-side part of the Sybase replication manager.
- Data Modelling for data replications.
- Modeling can be done in several places (bottom-up description).
- If data services is used to create and fill the table, first modeling decisions can be made here.
- Data models can be created within the In Memory Computing Engine.
- Models are stored in form of views and associated metadata in the engine.
- The front-end tool to create these models in the In-Memory Computing Studio (Information Modeler within that tool).
- Depending on the front-end tool used to retrieve data from the In-Memory Computing Engine, further modeling decisions can be made in universes (SAP Business Objects Information Design Tool) or other semantic layers.

Architecture Overview

Persistence Layer in In-memory Comp. Engine

System Restart

Reboot or Power failure deletes in-memory data

- ➔ System is normally restarted („lazy" restart to keep downtime short: tables with preload flag + subsequently requested tables are loaded first)
- ➔ System is restored to the state just before the failure (except non-committed transactions)
- Used for recovery:
  - Last data savepoint
  - Log between the last data savepoint and the time of failure (contains the data changes of all commited transactions up to that point)

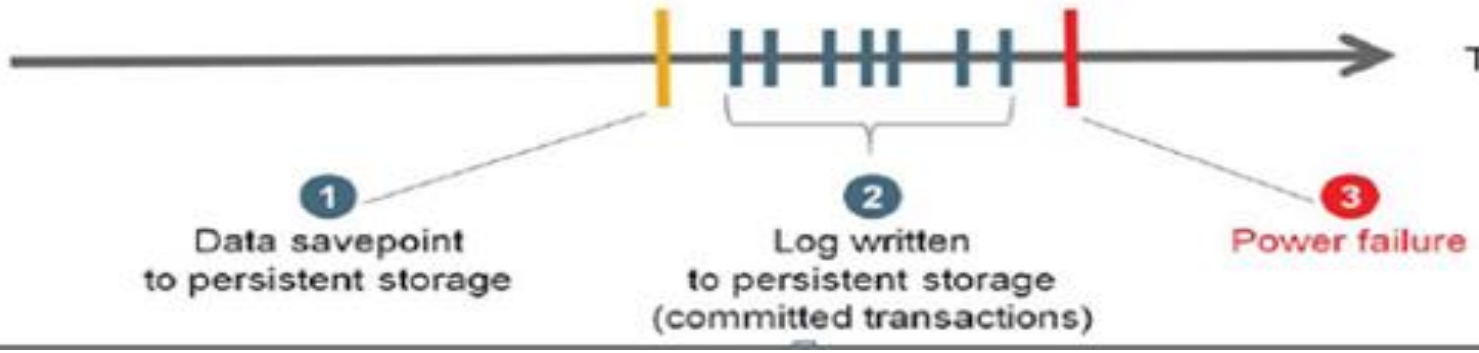

In Memory Design

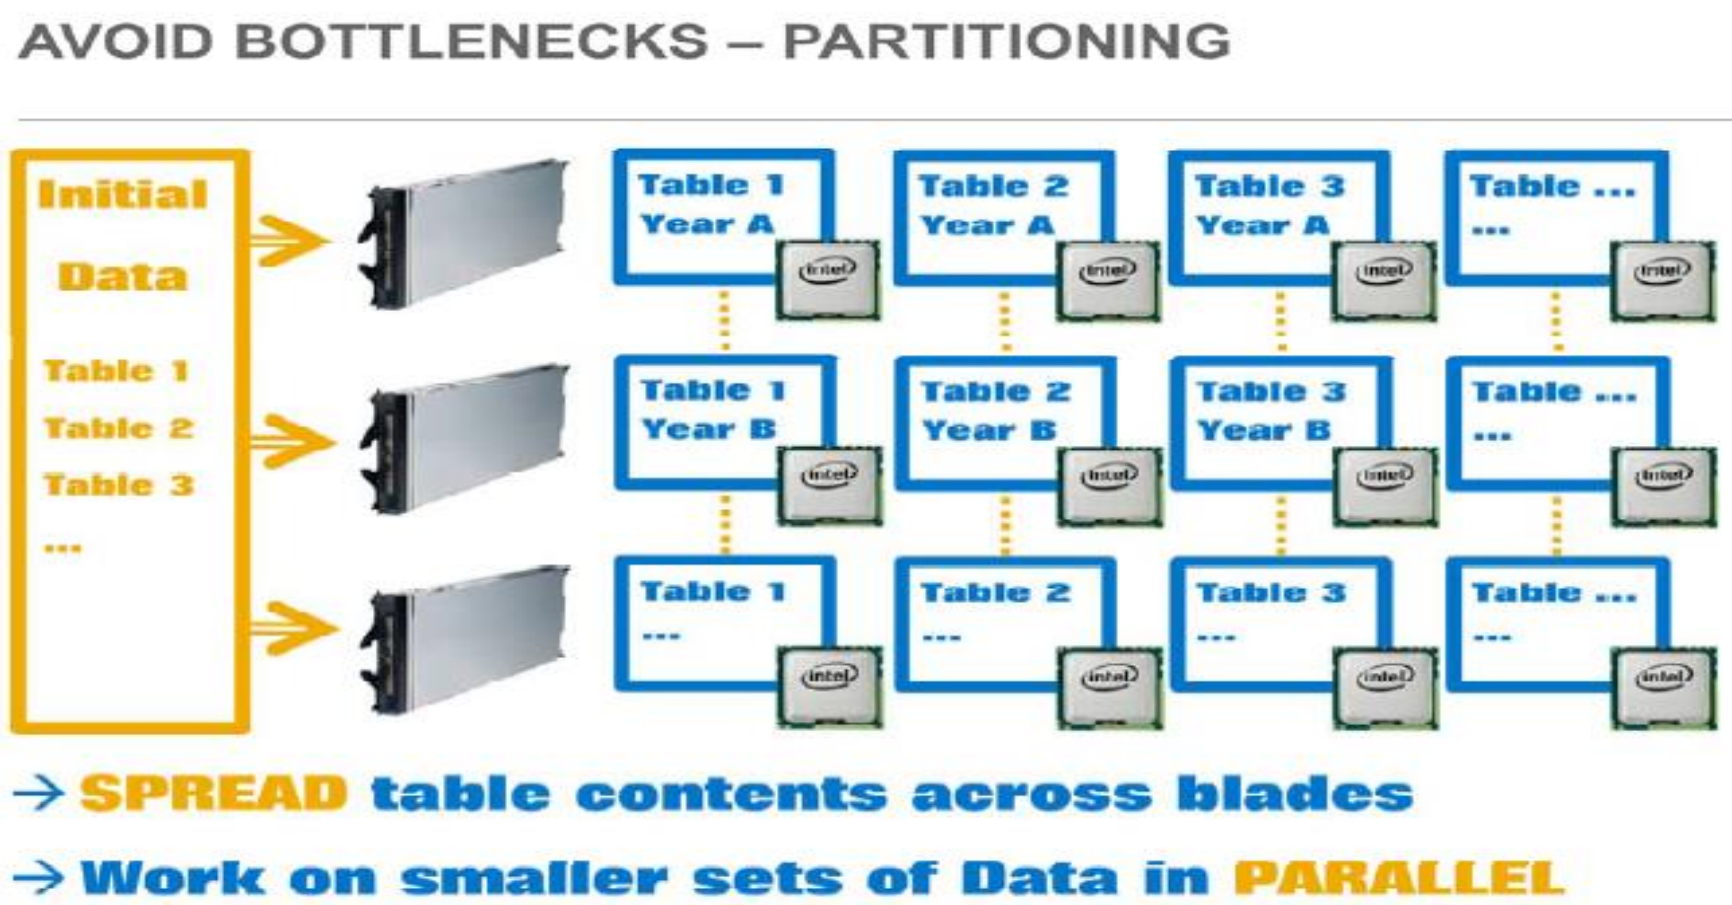

Parallel Processing

SAP HANA was designed to perform its basic calculations, such as analytic joins, scans and aggregations in parallel. Often it uses hundreds of cores at the same time, fully utilizing the available computing resources of distributed systems.

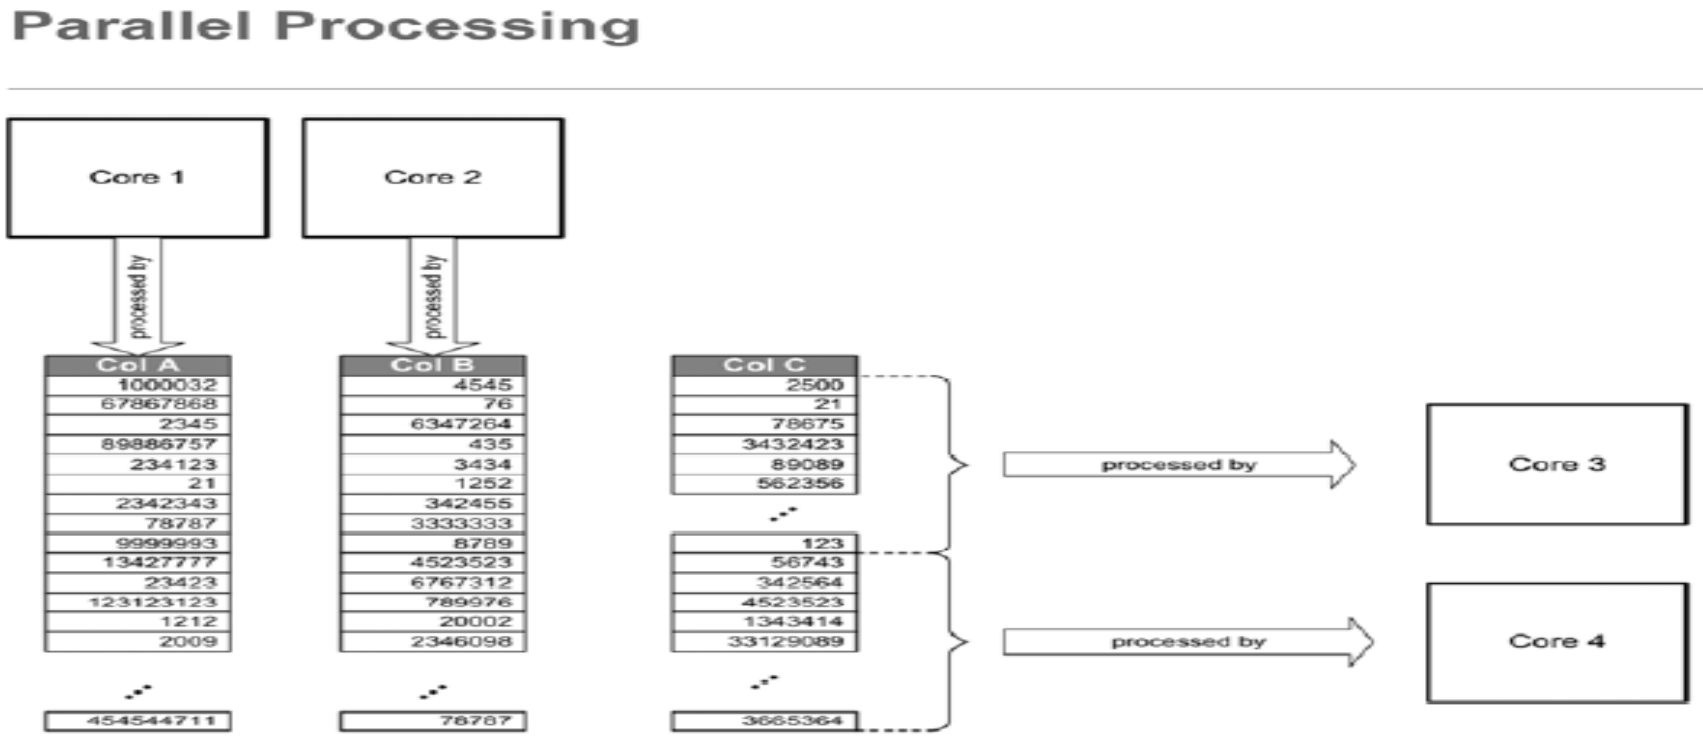

Row Store and Column Store

A database table is conceptually a two-dimensional data structure organized in rows and columns. Computer memory, in contrast, is organized as a linear structure. A table can be represented in row-order or column-order.

Table

Country	Product	Sales
US	Alpha	3.000
US	Beta	1.250
JP	Alpha	700
UK	Alpha	450

Row Store

Row 1	US
	Alpha
	3.000
Row 2	US
	Beta
	1.250
Row 3	JP
	Alpha
	700
Row 4	UK
	Alpha
	450

Column Store

Country	US
	US
	JP
	UK
Product	Alpha
	Beta
	Alpha
	Alpha
Sales	3.000
	1.250
	700
	450

Row Store

Row Store

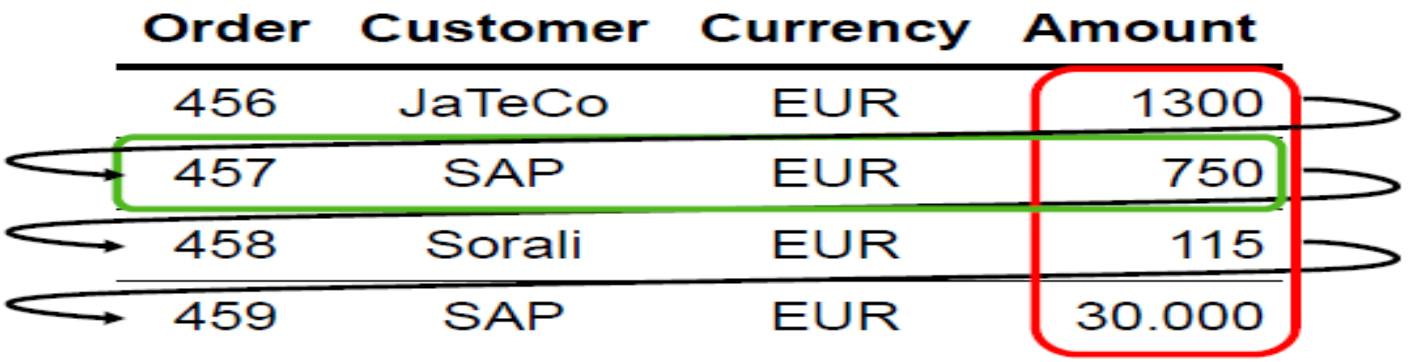


```
SELECT * ….
 WHERE ORDER = 457
```

Good performance

```
SELECT SUM(Amount)…
```

Low performance

Column Store

Column Store

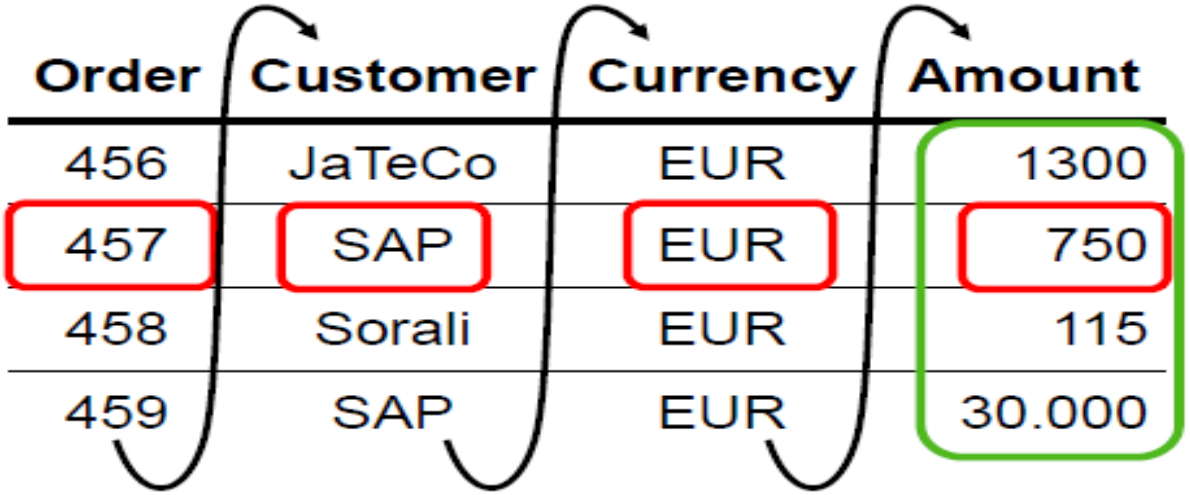

Order	Customer	Currency	Amount
456	JaTeCo	EUR	1300
457	SAP	EUR	750
458	Sorali	EUR	115
459	SAP	EUR	30.000

```
SELECT * ….
 WHERE ORDER = 457
```

Low performance

```
SELECT SUM(Amount)…
```

Good performance

When to use row storage

When to Use Row Storage?

Row storage is more suitable for tables

- that contain mainly distinct values
  → low compression rate
- in which most/all columns are relevant
- that are not subject to aggregation or search operations on non-indexed columns
- that are fully buffered
- that have a small number of records

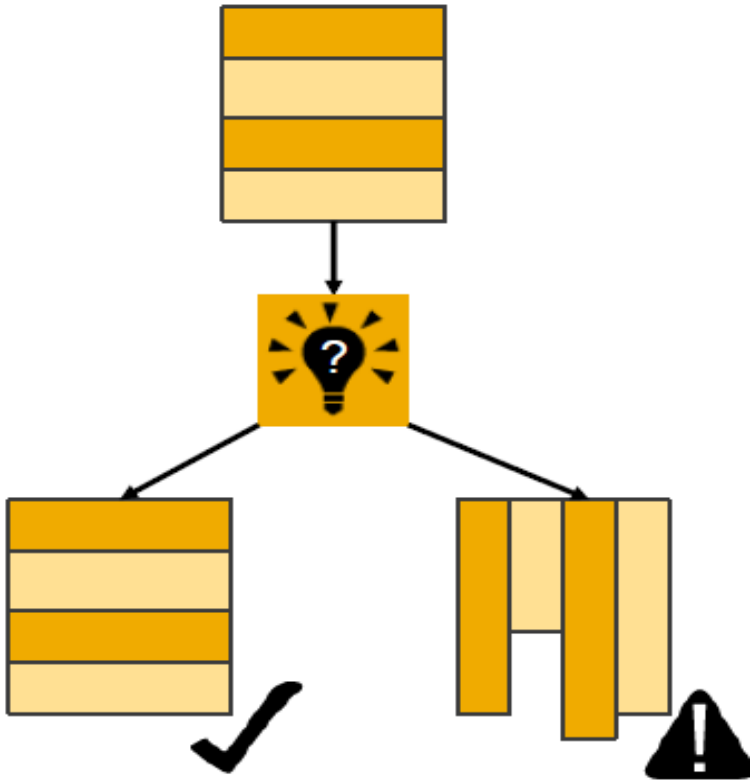

When to use columnar storage

When to Use Columnar Storage?

Columnar storage is best for tables

- that are subject to column operations on a large number of rows
- that have a large number of columns, more unused
- that are subject to aggregations and intensive search operations

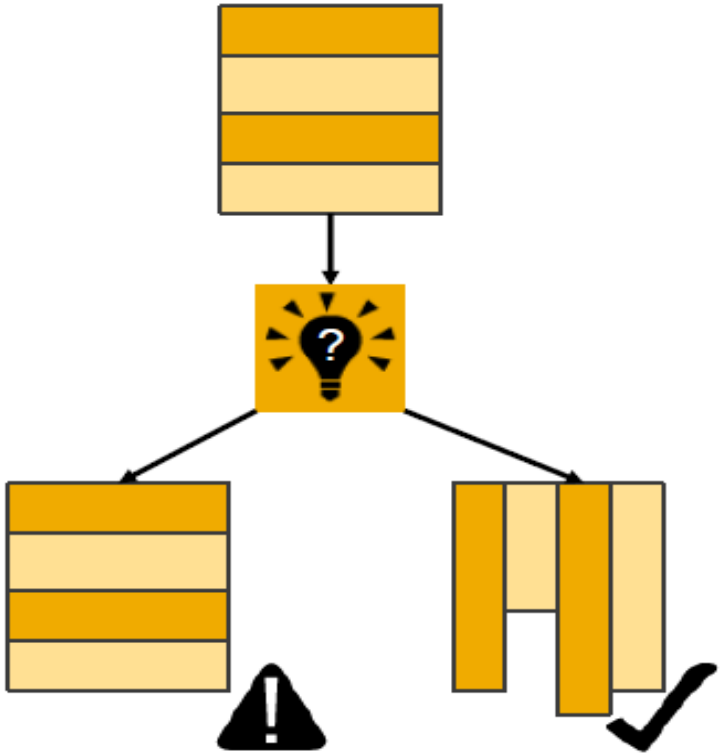

Dictionary Compression

Column Store Dictionary Compression

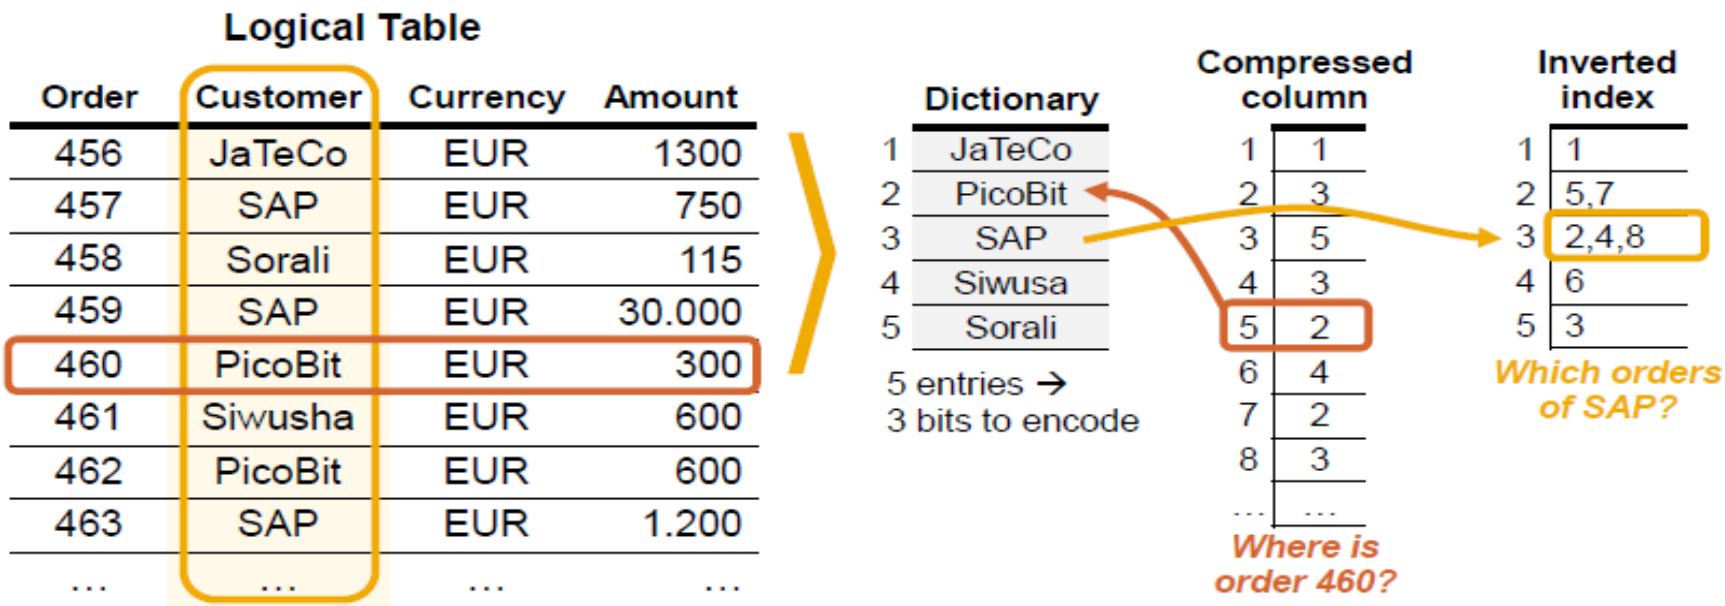

SAP HANA Coding Pattern

In-memory computing imperative:

- Avoid (unnecessary) movement of large data volume
- Perform data-intensive calculations in the database

Today:

Data-intensive computations in application layer

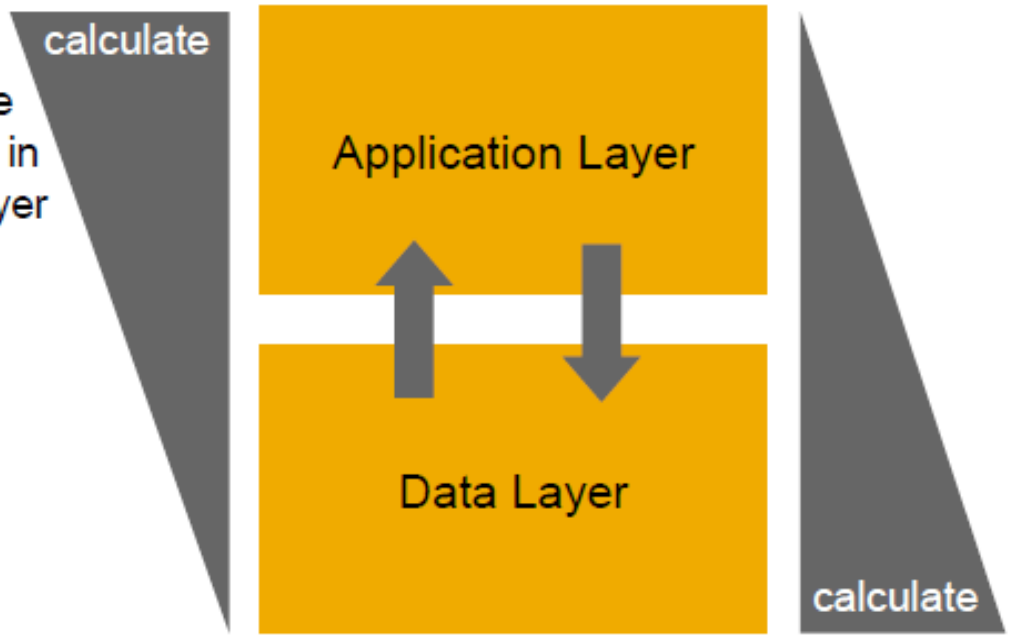

SAP HANA:

Delegate data-intensive operations to data layer

Introduction to SAP HANA Studio

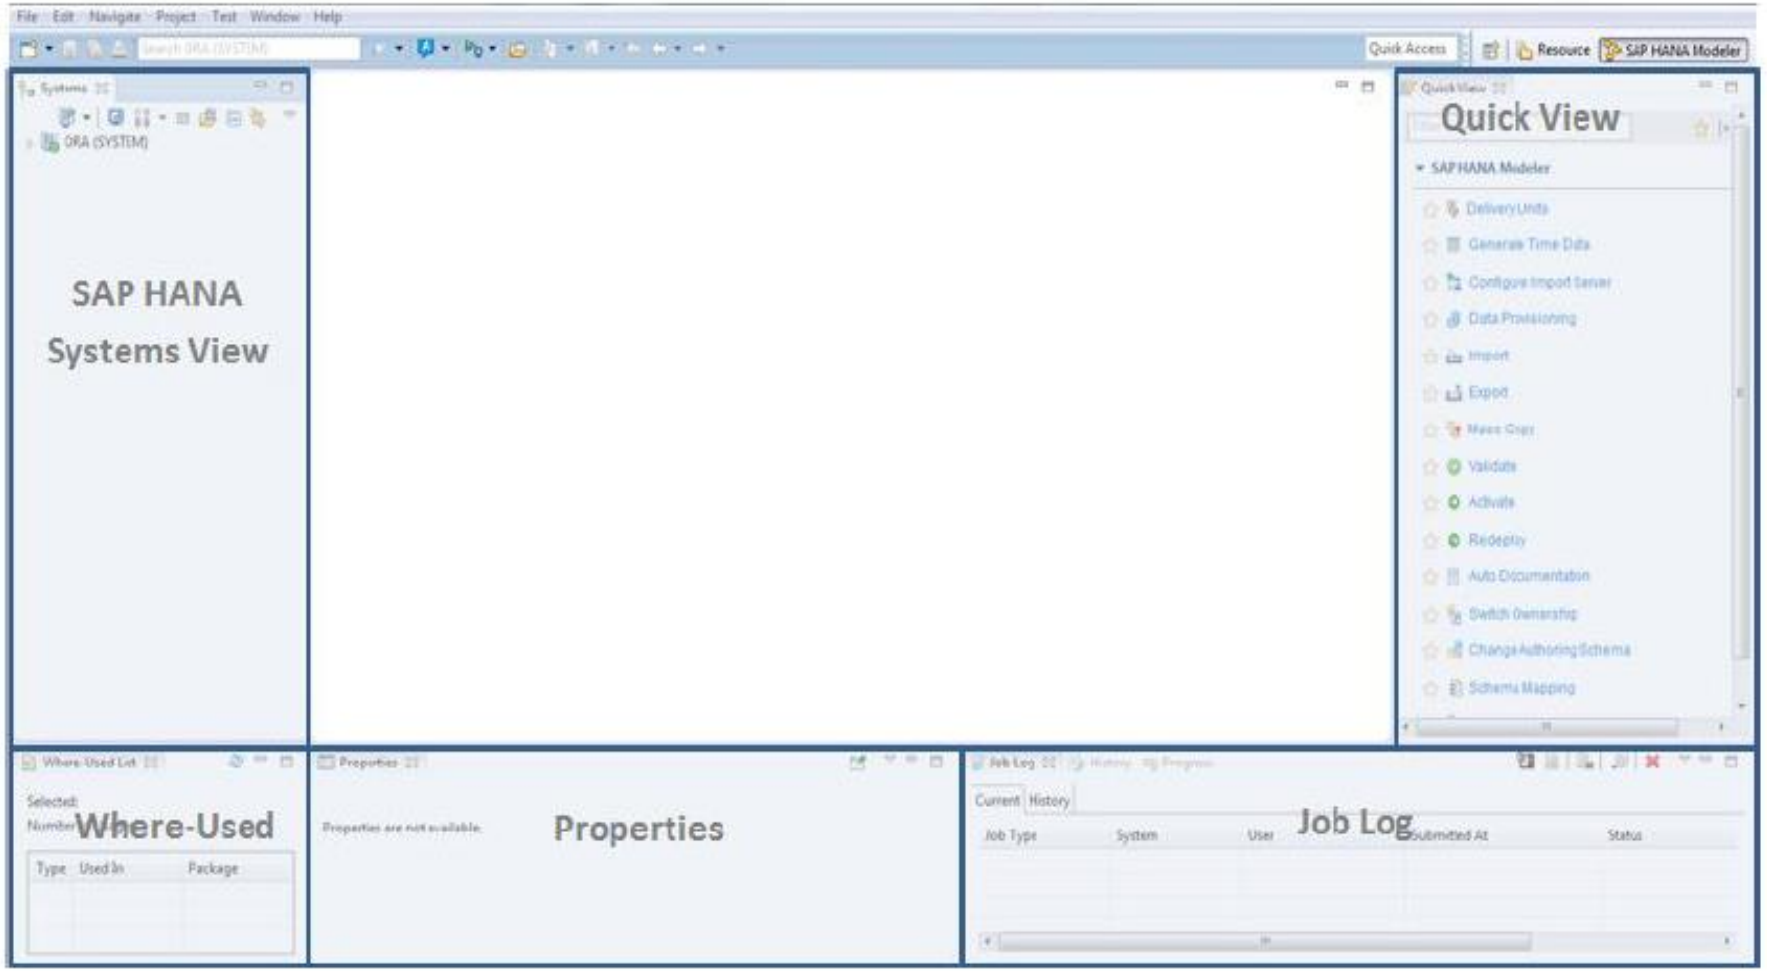

Administration Console

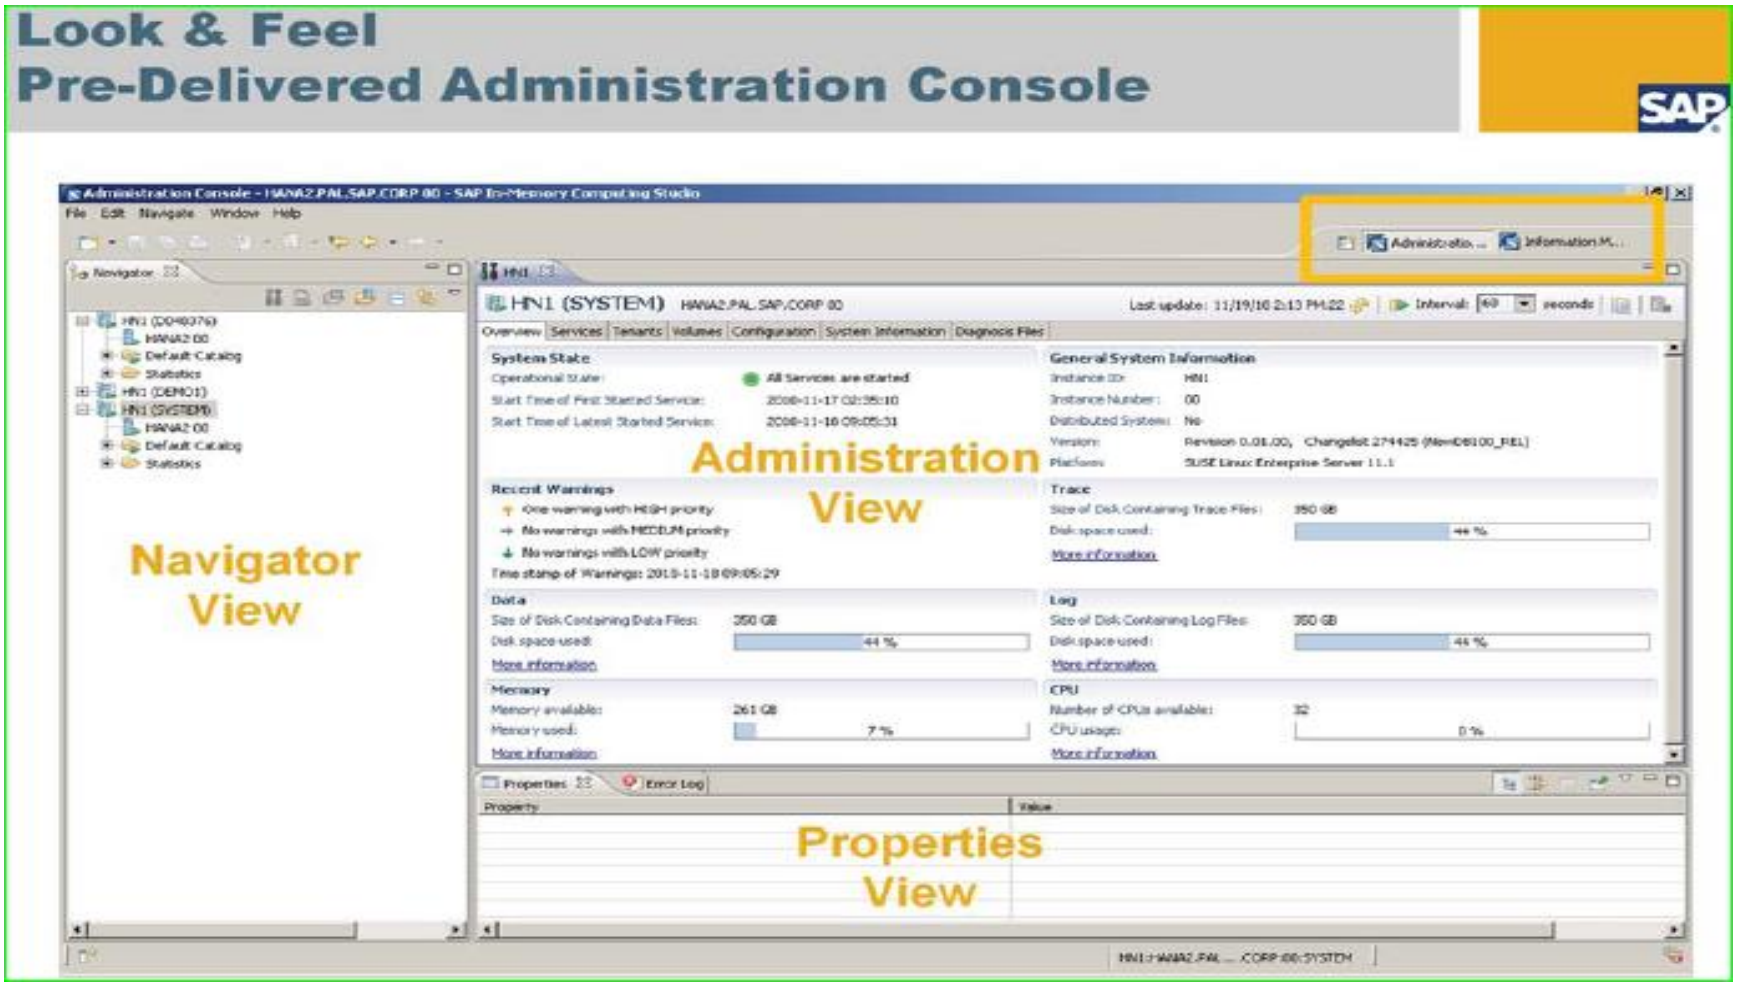

Introduction to SAP HANA Studio

SAP HANA studio: Editing Tools

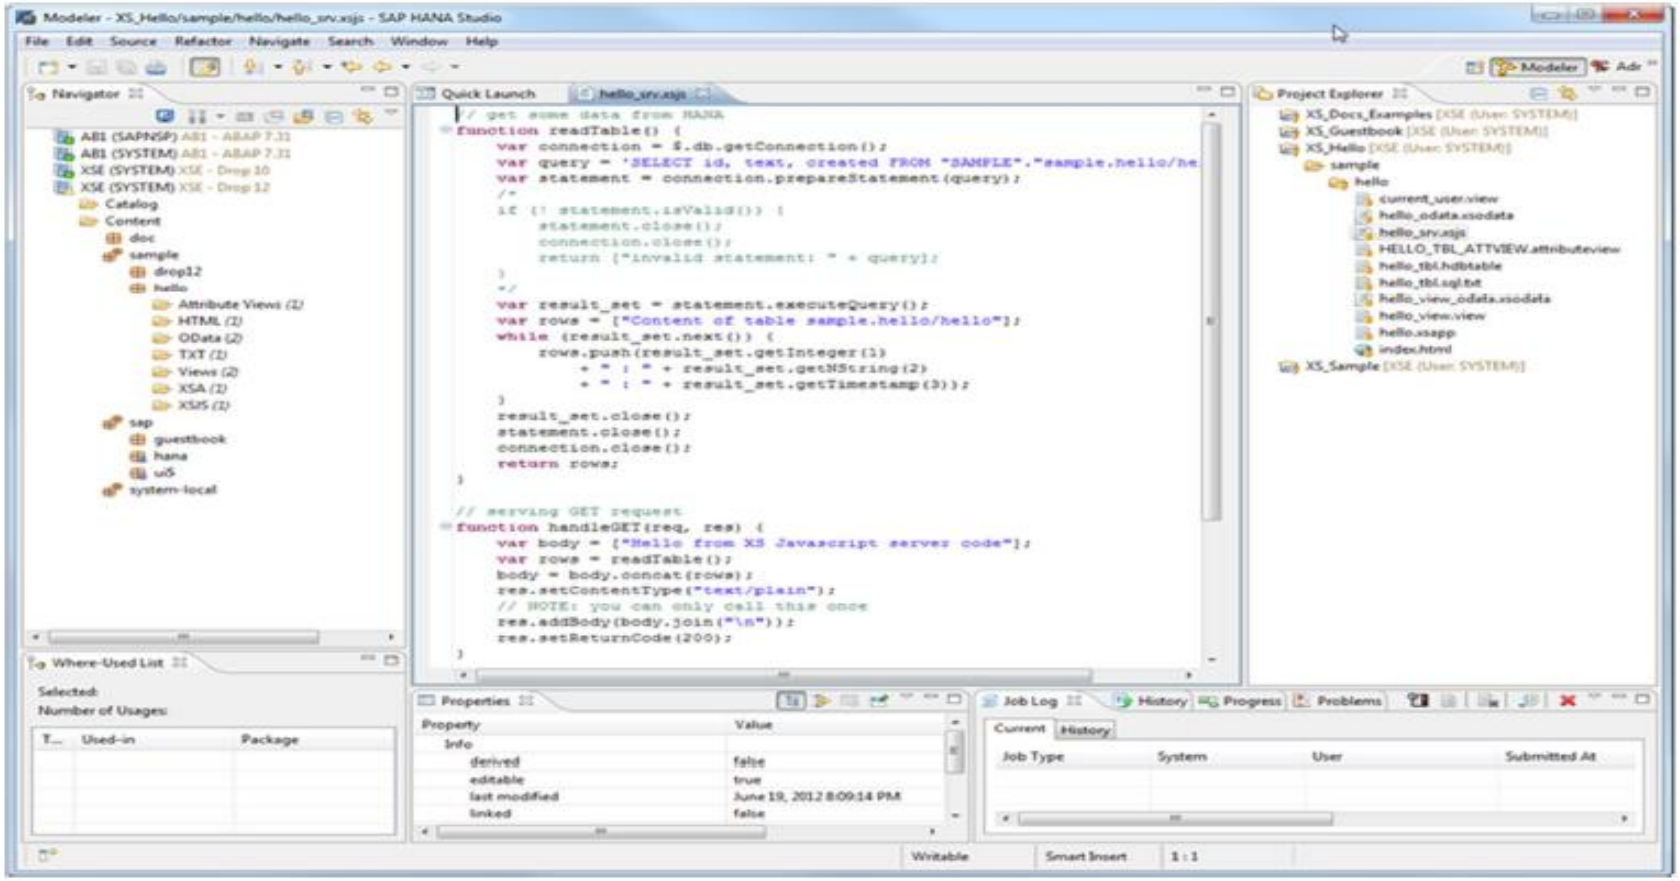

ABAP Code on HANA Studio

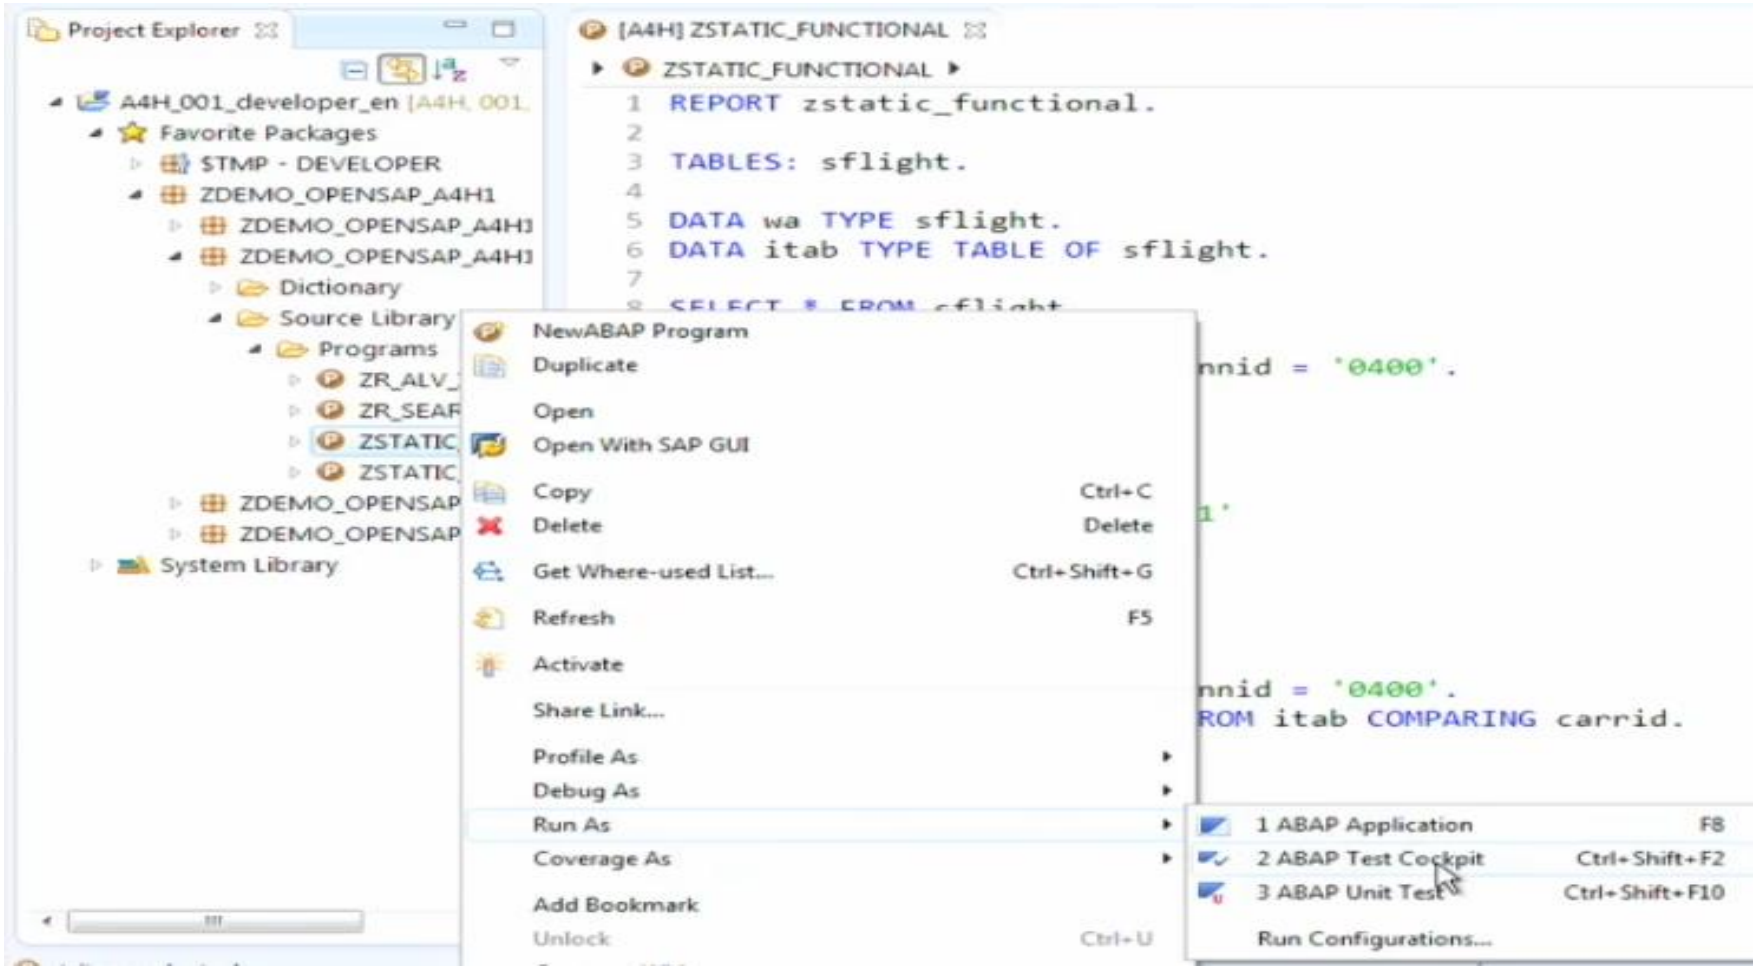

ABAP Code and Errors on HANA Studio

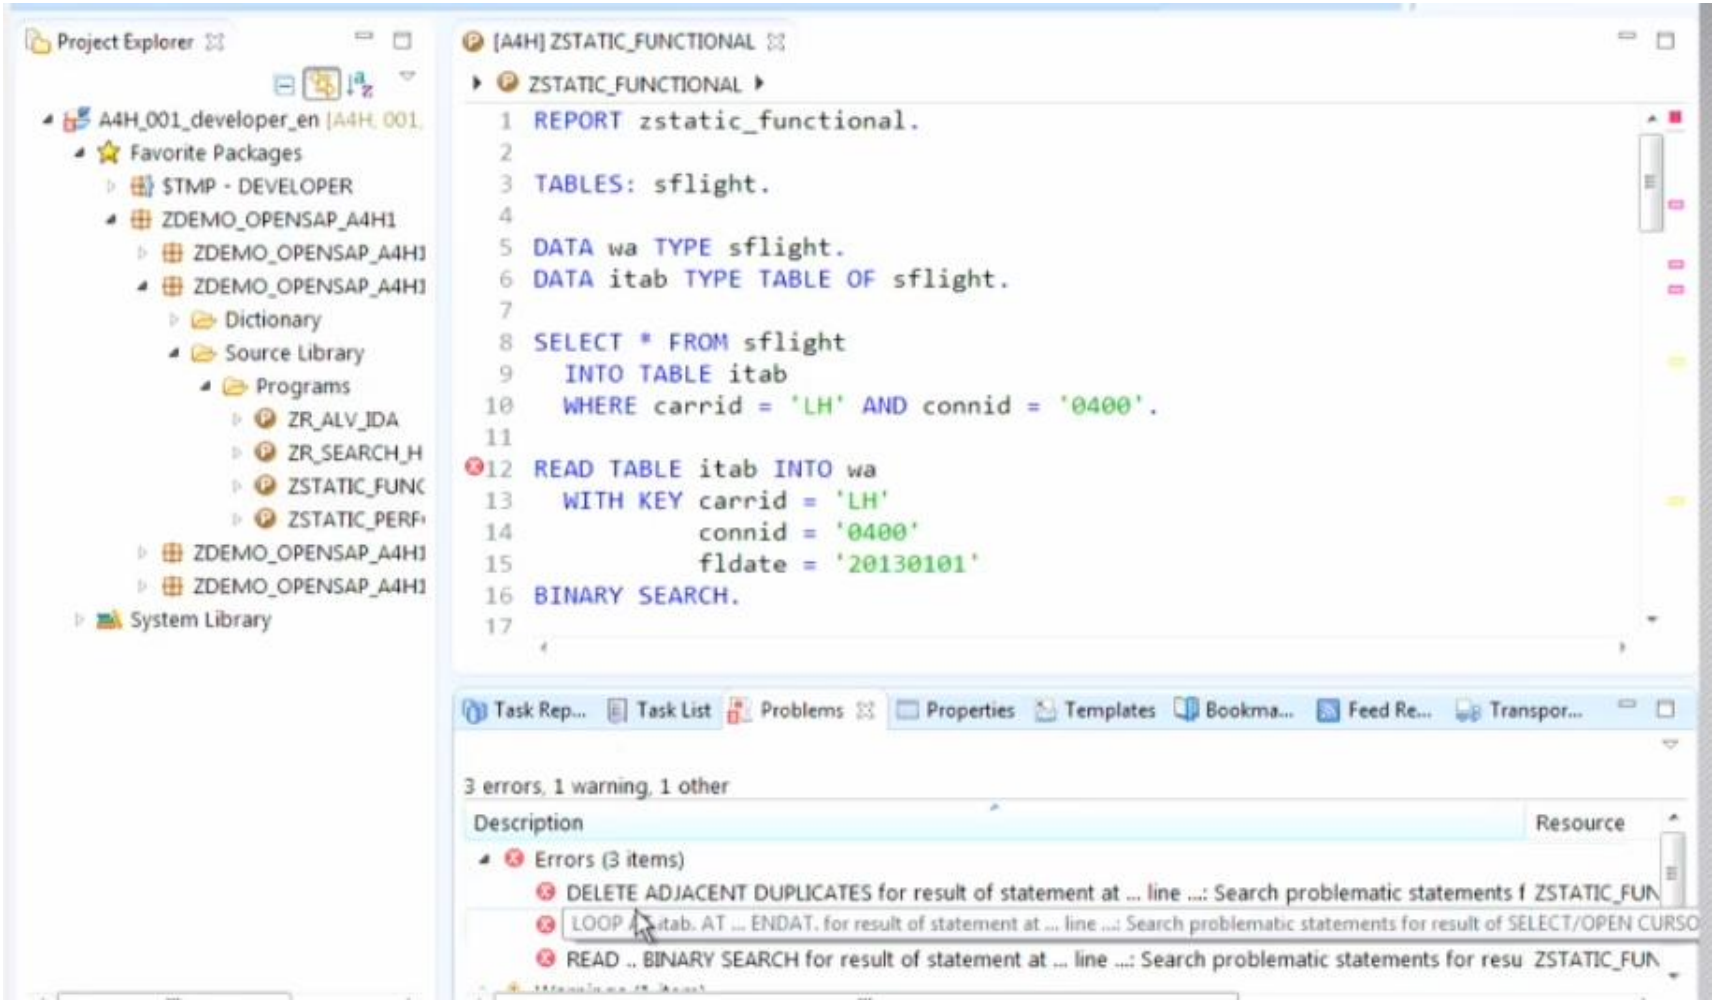

ABAP Code and Help on HANA Studio

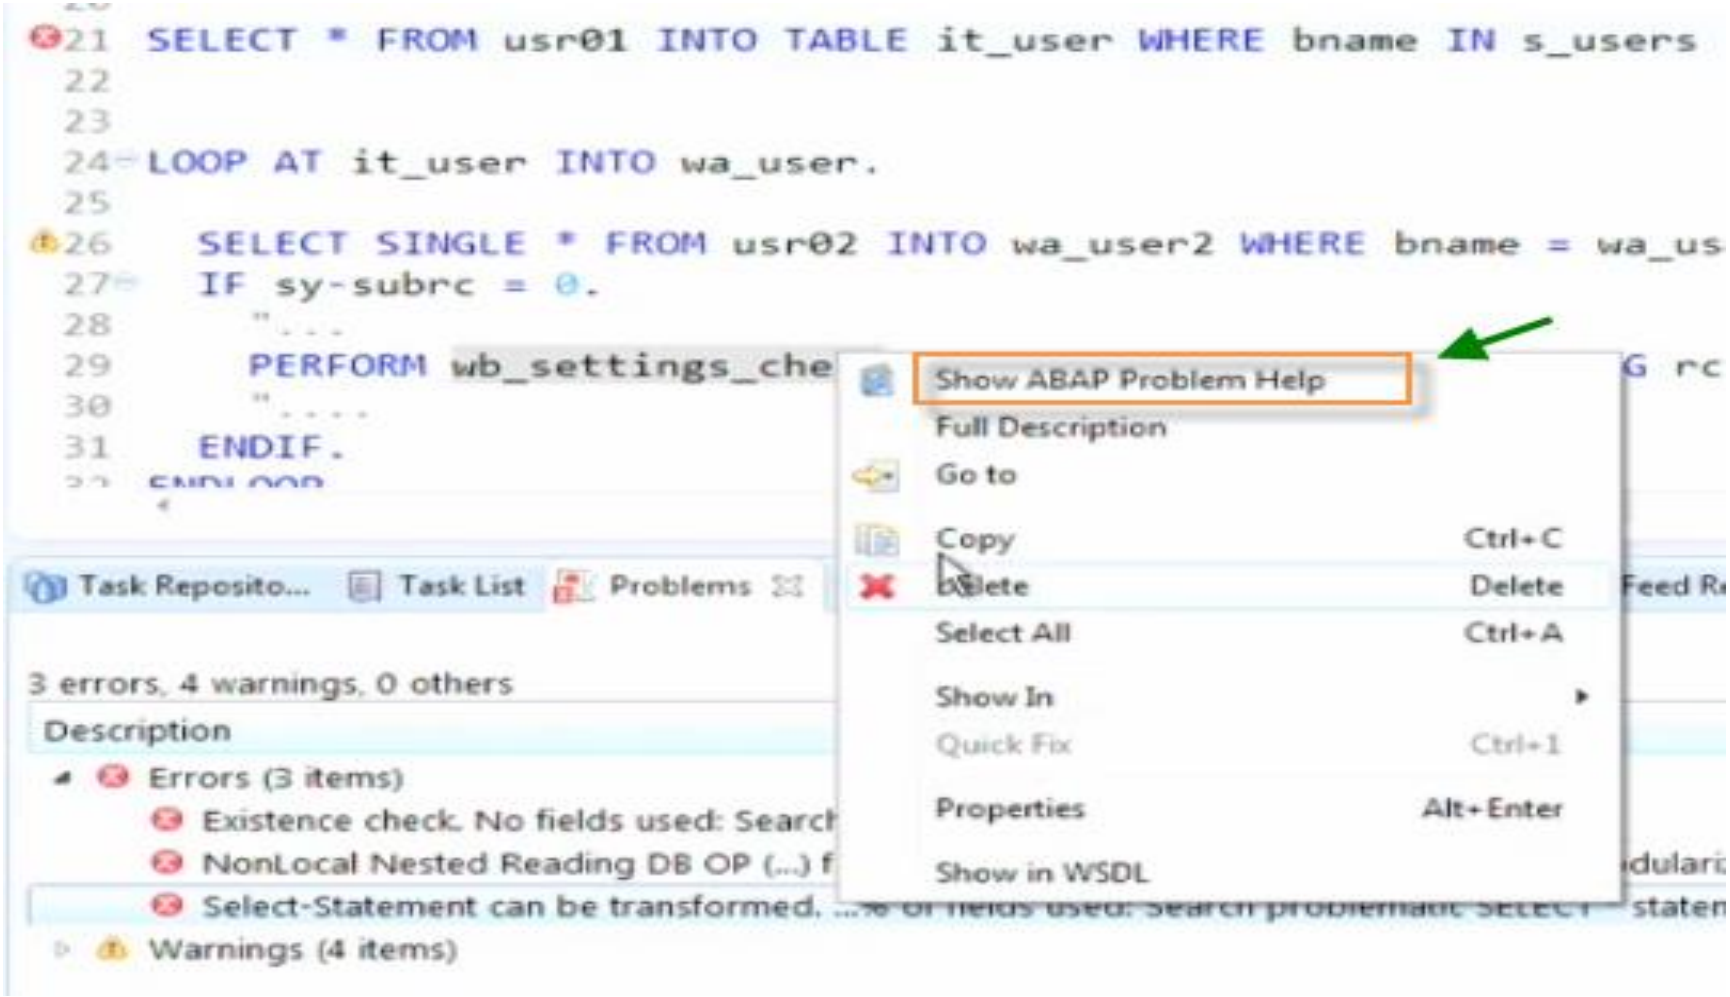

ABAP Analysis on HANA Studio

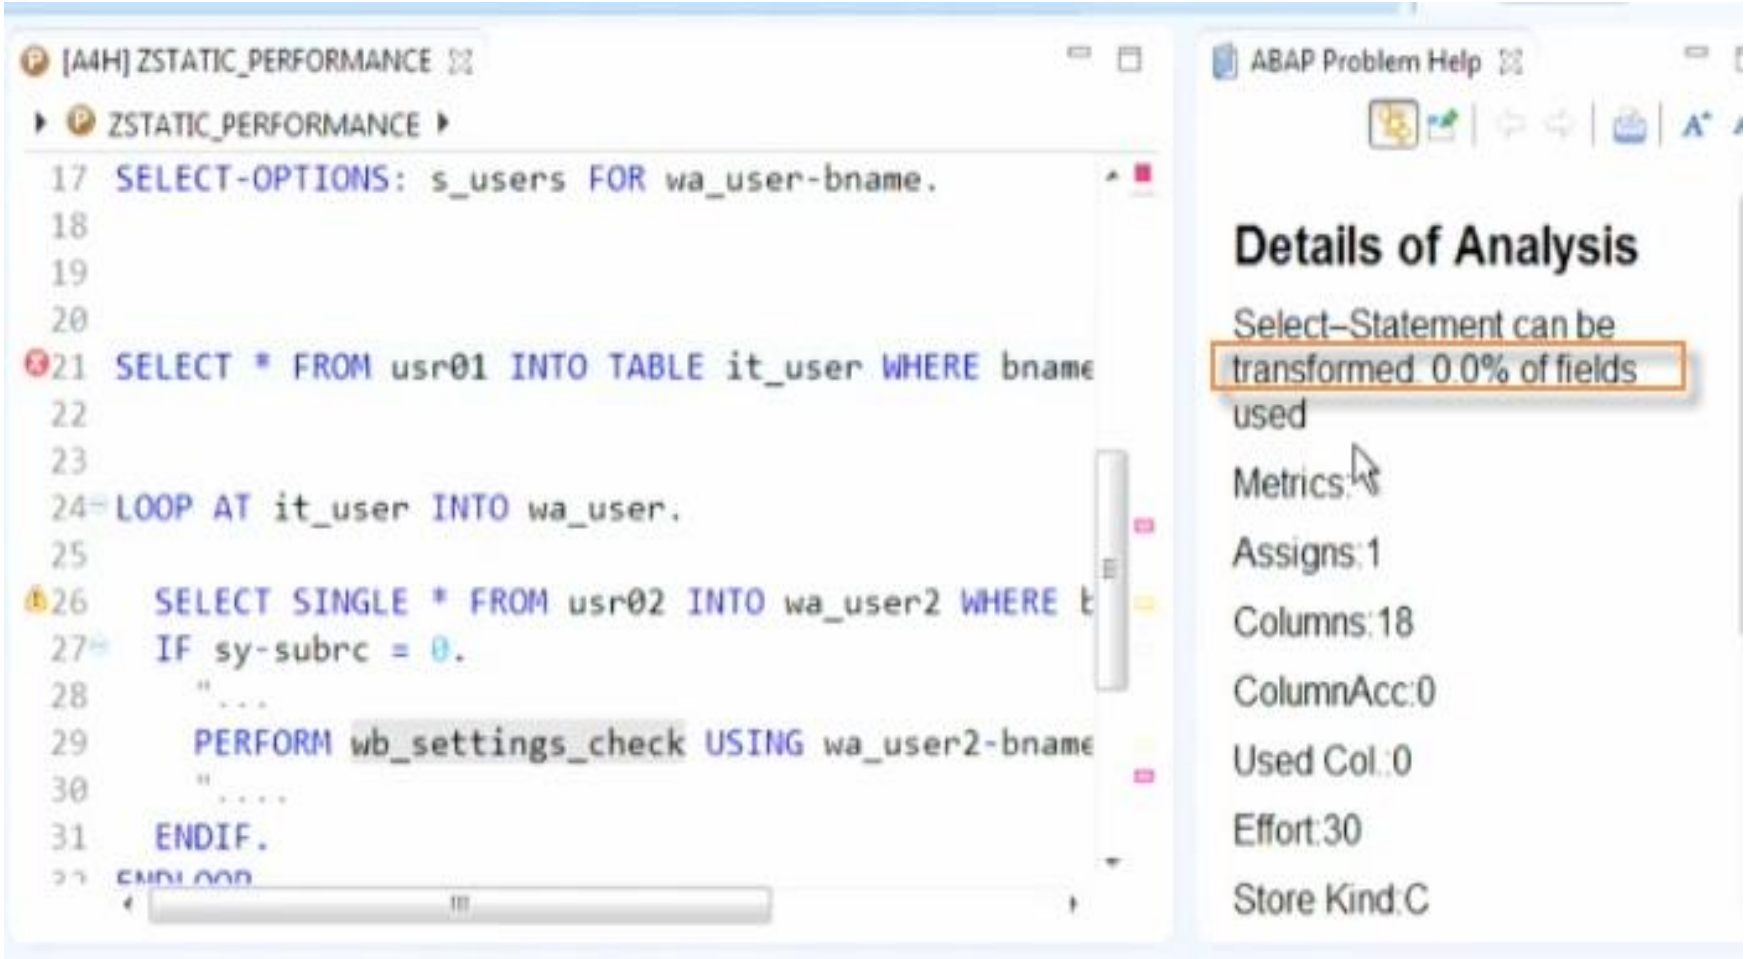

SAP HANA Mandatory Adaptations

Mandatory Adaptations: Examples

Previous usage of Native SQL/hints

- Check if Native SQL uses vendor-specific statements and how they can be replaced by either Open SQL or native SAP HANA constructs.

```
" Use of native SQL
EXEC SQL.

ENDEXEC.

" Native SQL via ADBC
lo_result =
  NEW cl_sql_statement(
   )->execute_query(
'SELECT ROWNUM, *
   && '   FROM SNWD_BPA'
   && ' ORDER BY COMPANY_NAME' ).
```

- Check if previously used hints are still needed in adapted form or can be removed.

```
" Database / Database Interface Hints
SELECT * FROM snwd_so
  INTO wa FOR ALL ENTRIES IN fae
 WHERE node_key = fae-node_key
  %_HINTS MSSQLNT '&prefer_join 0&'.
```

SAP HANA Mandatory Adaptations

Mandatory Adaptations: Examples

Relying on undocumented behavior

- Relying on implicit sorting

```
" Relying on implicit sorting
SELECT * FROM my_table
  INTO TABLE lt_data  WHERE id < 100.
```

```
READ TABLE lt_data
  WITH KEY id = 10 BINARY SEARCH.
```

- Direct access to physical pool/clusters

```
" Access to physical pool / cluster
DELETE FROM my_cluster
 WHERE timestamp <  '01012000'.
```

- Checking for existence of secondary indices

```
" Check for secondary index
IF ( lv_index_exists = abap_true ).
  ...
ENDIF.
```

ABAP Development Shortcuts

Speed up your ABAP development using shortcuts

Edit	
Ctrl+Shift+A	Open development object
Ctrl+F2	Check development object
Ctrl+F3	Activate development object
Ctrl+Shift+F3	Activate all inactive objects
Ctrl+Space	Code completion
Ctrl+1	Quick fix proposal
Ctrl+<	Add comment
Ctrl+Shift+<	Remove comment
Shift+F1	Format source aka pretty printer
Help	
F1	ABAP keyword documentation
F2	Show code element information
Ctrl+3	Search for commands & views
Ctrl+Shift+L	List all keyboard shortcuts

Navigate	
F3	Open definition
Alt+Left	Backward history
Alt+Right	Forward history
Ctrl+T	Quick hierarchy
F4	Open Type Hierarchy
Ctrl+O	Quick outline
Ctrl+Shift+G	Where-used list
Run, Debug	
F8	Run current ABAP object
Alt+F8	Select & run ABAP application
Ctrl+Shift+B	Toggle breakpoint
F5, F6, F7, F8	Step into, over, return, resume
Ctrl+Shift+F10	Execute ABAP unit tests
Alt+F9	Profile development object

ABAP Development Shortcuts

Shortcut	What it will do
CTRL+D	Deletes the selected codeline
CTRL+SHIFT+DELETE	Deletes the content from the cursor postion to the end of the line
CTRL+DELETE	Deletes the next word in the editor
CTRL+BACKSPACE	Deletes the previous word in the editor
ALT+UP/DOWN	Moves the selected codelines up and down in the editor
CTRL+ALT+UP/DOWN	Duplicates Codelines before/after the selected codeline
CTRL+UP/DOWN	Scrolls Line up and down
SHIFT+ENTER	Adds a new line below the current line and positions the cursor in that line
CTRL+SHIFT+ENTER	Adds a new line above the current line and positions the cursor in that line
CTRL+Z	Undo changes
ALT+SHIFT+R	Renames the selected object, e.g. variable, method, class
CTRL+1	Opens Quickfix/Quickassist Dialog on the selected element
CTRL+7	Comments/Uncomments selected code in the editor
SHIFT+F1	Formats the source code (aka. Pretty Printer)
CTRL+N	Creates new development object

ABAP Development Shortcuts

Shortcuts for the Editor Tabs:

Shortcut	What it will do
CTRL+E	Displays a list of all open editors
CTRL+F6	Easily switch between the editor tabs (Like Tab for Windows)
CTRL+F7	Easily switch between all eclipse views
CTRL+F8	Easily switch between the perspectives
CTRL+M	Maximize the active editor or viewer to full-screen mode
CTRL+3	Easily open Eclipse views or trigger command via the Quick Access Inputfield
CTRL+PAGE UP/PAGE DOWN	Navigate through the editor tabs forward and backward
F3	Open new editor tab based on the cursor position
ALT+PAGE UP/PAGE DOWN	Navigate through the tabs of the class editor between global class, local class and test classes
CTRL+F4	Close the active editor tab
CTRL+SHIFT+F4	Close all editor tabs

ABAP Development Shortcuts

Shortcuts for Navigation:

Shortcut	What it will do
CTRL-L	Jump to line in editor
CTRL-O	Launch the Quick Outline to easily navigate to methods, attributes etc.
ALT+LEFT/RIGHT	Navigate through the editor navigation history
CTRL+; / :	Step quickly through the editor markers, like tasks, bookmarks, error markers, ATC findings etc.
F3	Navigate to the definition of the selected element, e.g. variable, method, attribute etc.

Summary

- By end of this course, participants know
  - Understand SAP HANA
  - To work with ABAP ON HANA
  - Understand SAP HANA System Architecture
  - To know SAP IN-Memory Strategy and Technology
  - To work with Row store and Column Store tables

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Thank you